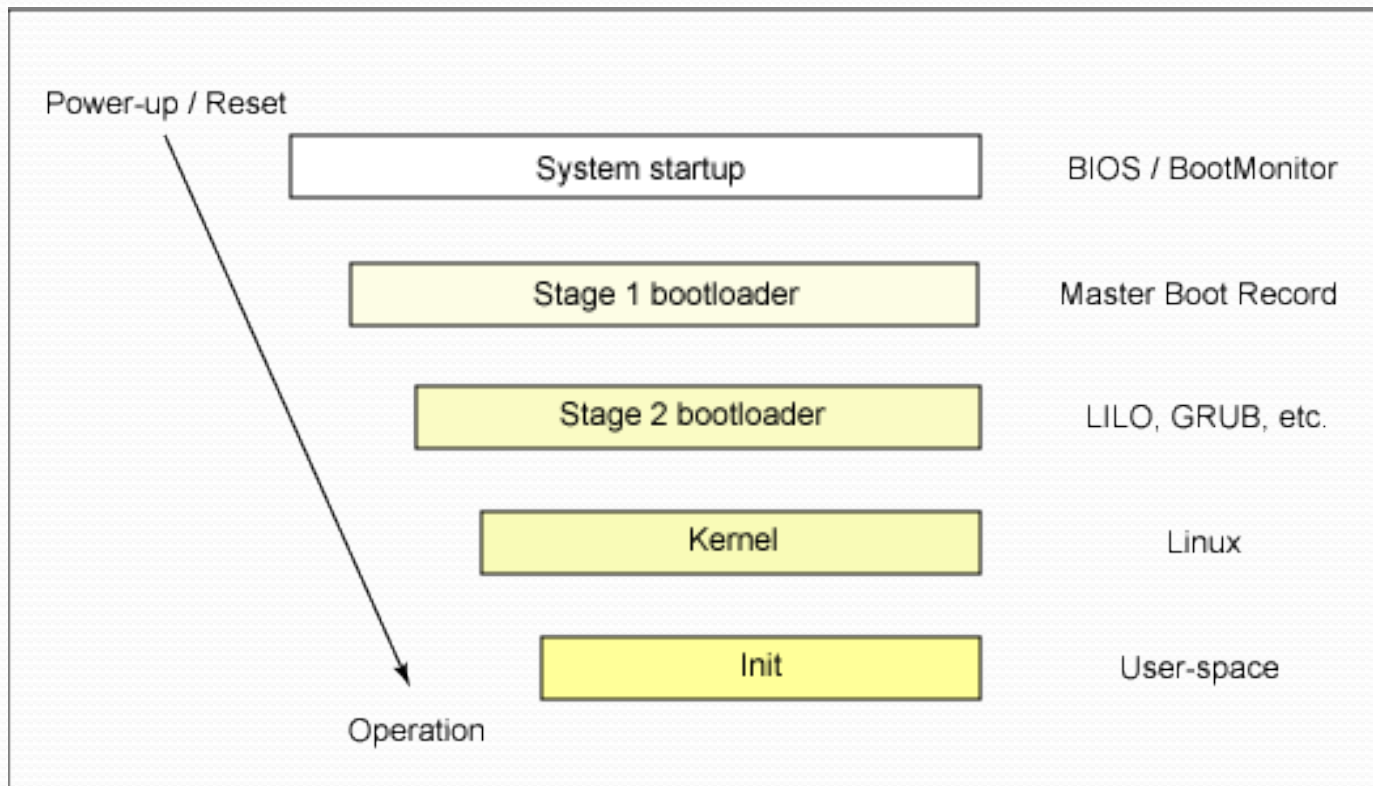


Boot Process

Presented By:

Amir Reza Niakan Lahiji

Boot Process Overview



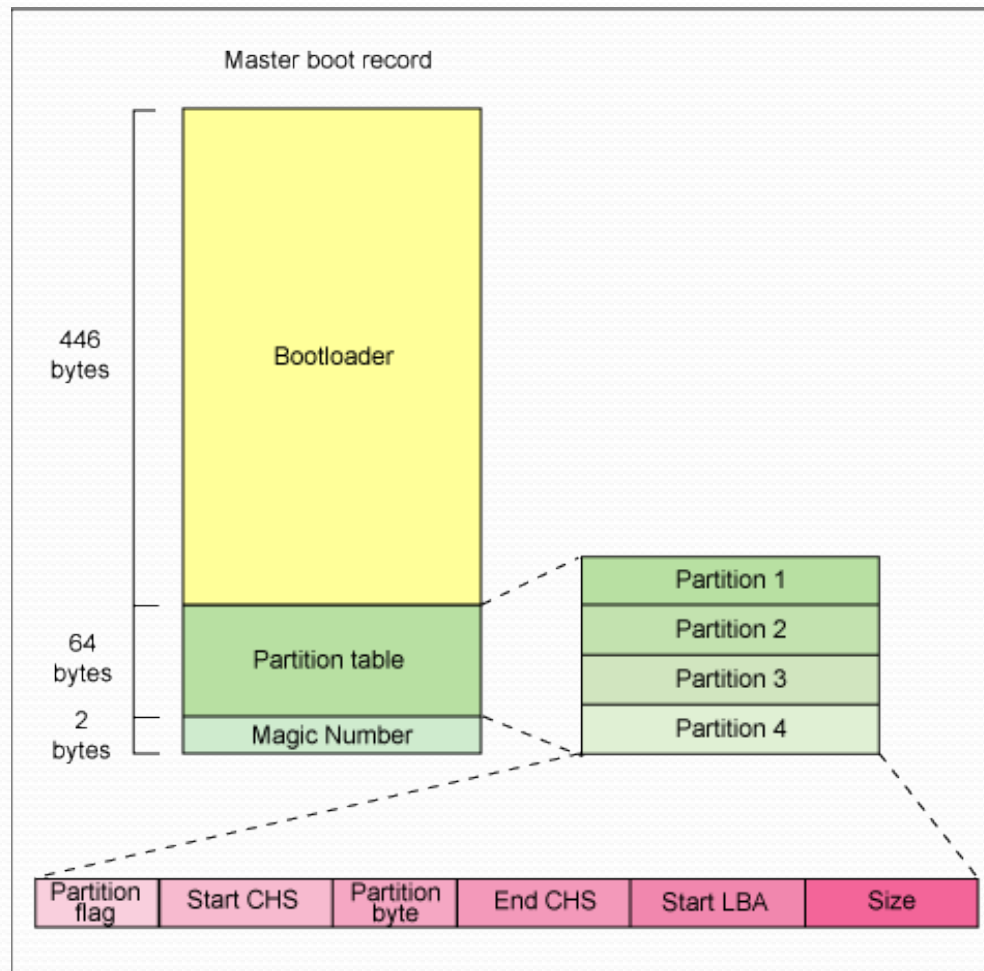
System Startup in PC

- booting Linux begins in the BIOS at address 0xFFFF0
 - First step: Power-On Self Test (POST)
 - Second step: local device enumeration and initialization
 - Searches for devices that are both active and bootable based on the preferred order which is determined in CMOS
 - If hard disk must be boot(Suppose Linux resides in hard disk)
 - BIOS loads MBR (Master Boot Record) in RAM and yields control to it
 - MBR contains the primary boot loader
 - MBR is a 512-byte sector, located in the first sector on the disk (sector 1 of cylinder 0, head 0)

Stage 1 boot loader

- The primary boot loader is a 512-byte image which contains
 - Program Code (first 446 bytes)
 - Partition Table (64 bytes)
 - Magic Number (2 bytes)
 - The job of the primary boot loader is to find and load the secondary boot loader
 - First looks through the partition table for an active partition
 - When it finds an active partition, it scans the remaining partitions in the table to ensure that they're all inactive
 - When this is verified, the active partition's boot record is read from the device into RAM and executed

Stage 1 boot loader



Stage 2 boot loader (Kernel loader)

- The task at this stage is to load the Linux kernel and optional initial RAM disk.
- The first- and second-stage boot loaders combined are called Linux Loader (LILO) or GRand Unified Bootloader (GRUB) in the x86 PC environment.
- The great thing about GRUB is that it includes knowledge of Linux file systems.
 - GRUB can load a Linux kernel from an ext2 or ext3 (and also ext4) file system
 - It does this by making the two-stage boot loader into a three-stage boot loader.
 - Stage 1 (MBR) boots a stage 1.5 boot loader that understands the particular file system containing the Linux kernel image

Kernel loader

- Kernel image is compressed and typically this is
 - a zImage (compressed image, less than 512KB) or
 - a bzImage (big compressed image, greater than 512KB)
- At the head of this kernel image is a routine that does some minimal amount of hardware setup and then decompresses the kernel contained within the kernel image and places it into high memory.
- If an initial RAM disk image is present, this routine moves it into memory and notes it for later use.
- The routine then calls the kernel and the kernel boot begins
- During the boot of the kernel, the initial-RAM disk (initrd) that was loaded into memory by the stage 2 boot loader is copied into RAM and mounted.

Kernel loader

- This initrd serves as a temporary root file system in RAM and allows the kernel to fully boot without having to mount any physical disks
- initrd is mainly designed to allow system startup to occur in two phases, where the kernel comes up with a minimum set of compiled-in drivers, and where additional modules are loaded from initrd.
- After the kernel is booted, the root file system is pivoted where the initrd root file system is unmounted and the real root file system is mounted.
- The initrd function provides a means of bootstrapping to gain access to the disk and mount the real root file system.

Kernel loader

- Let see
 - `ls -l /boot`
 - `vmlinuz-2.6.28-11-generic` → compressed kernel image
 - `initrd.img-2.6.28-11-generic` → initial RAM disk
 - `ls -l /boot/grub`
 - `stage1` -> 512 byte -> primary boot loader
 - `e2fs_stage1_5` -> ex2, ex3 or ext4 file systems
 - `fat_stage1_5` -> fat file systems
 - ...
 - `stage2`

Init

- After the kernel is booted and initialized, the kernel starts the init (first user-space application).
- System V init process
 - Run levels concept
 - /etc/inittab
- Upstart init process
 - Event-driven init process
 - Ubuntu 6.10 and later
 - Fedora 9.0 and later

System V init

- Runlevel
 - is a number which indicates what "mode" you want to computer to boot into
 - 0 — Halt
 - 1 — Single-user mode
 - 2 — Not used (user-definable)
 - 3 — Full multi-user mode
 - 4 — Not used (user-definable)
 - 5 — Full multi-user mode (with an X-based login screen)
 - 6 — Reboot

System V init

- The default runlevel for a system to boot to and stop is configured in `/etc/inittab`. Like:
 - `id:3:initdefault:`
- The `/etc/init.d` directory contains the scripts executed by `init` at boot time and when the `init` state (or "runlevel") is changed
- These scripts are referenced by symbolic links in the `/etc/rcn.d` directories
- The names of the links all have the form *Smmscript* or *Kmmscript* where *mm* is a two-digit number and *script* is the name of the script (this should be the same as the name of the actual script in `/etc/init.d`).
- When changing runlevels, `init` looks in the directory `/etc/rcn.d` for the scripts it should execute, where *n* is the runlevel that is being changed to, or `S` for the boot-up scripts.

System V init

- When init changes runlevel first the targets of the links whose names start with a K are executed, each with the single argument stop, followed by the scripts prefixed with an S, each with the single argument start.
- The two-digit number *mm* is used to determine the order in which to run the scripts: low-numbered links have their scripts run first.

System V init

- the chain of events for a SysV init boot is as follows:
 - The kernel looks in /sbin for init
 - init runs the /etc/rc.d/rc.sysinit script
 - rc.sysinit handles most of the boot loader's processes and then runs rc.serial (if it exists)
 - init runs all the scripts for the default runlevel
 - init runs /etc/rc.d/rc.local

Upstart init

- The Upstart init daemon is event-based and runs specified programs when something on the system changes
 - Instead of starting and stopping services only when the runlevel changes, Upstart can start and stop services upon receiving information that something on the system has changed.
- An *event* is a change in system state that init can be informed of.
 - the boot loader triggers the startup event
 - the system entering runlevel 2 triggers the runlevel 2 event
 - a filesystem being mounted triggers the path-mounted event
 - You can also trigger an event manually by using the `initctl emit` command.

Upstart init (jobs)

- A *job* is a series of instructions that init reads.
 - A *task* is a job that performs its work and returns to a waiting state when it is done
 - A *service* is a job that does not normally terminate by itself.
 - The /etc/event.d directory holds *job definition files* (files defining the jobs that the Upstart init daemon runs)
 - You can run and stop a job manually using the initctl start and stop commands, respectively.

Upstart init (task example)

- /etc/event.d/testtask

```
start on helloevent
script
    echo "hello world" > /home/test
    date >> /home/test
end script
```

- #initctl emit helloevent

Upstart init (service example)

- /etc/event.d/testtask

```
start on helloevent
```

```
respawn
```

```
exec /bin/service1 > output1.txt
```

- #initctl emit helloevent